

IX. Computing eigenvalues and eigenvectors

M. SOMBRA - 2024

Let A $n \times n$ -matrix and consider the algorithm:

QR iteration

$T_0 \leftarrow U_0^* A U_0$ for a given U_0 $n \times n$ unitary

for $k = 1, 2, \dots$

$T_{k-1} = Q_k \cdot R_k$ QR factorization

$T_k \leftarrow R_k \cdot Q_k$

Since $T_k = R_k \cdot Q_k = Q_k^* \cdot \overbrace{Q_k \cdot R_k}^{T_{k-1}} \cdot Q_k$

we have that

$$A = (Q_0 \dots Q_k) \cdot T_k \cdot (Q_0 \dots Q_k)^*$$

hence T_k is unitarily similar to A .

In most situations,

$$T_k \xrightarrow[k \rightarrow \infty]{} T \quad \begin{array}{l} \text{upper triangular} \\ \text{(Schur form)} \end{array}$$

This is less obvious!

To motivate this method and understand its convergence, we first study two other eigenvalue iterations: the power method and the orthogonal iteration.

IX.1 The power method [D, § 4.4.], [GVL, § 7.3.1]

It is the iteration:

Power method

Given $x_0 \in \mathbb{R}^n$ with $\|x_0\|_2 = 1$:
For $k=0, 1, 2, \dots$
 $y_{k+1} \leftarrow A \cdot x_k$
 $x_{k+1} \leftarrow \frac{y_{k+1}}{\|y_{k+1}\|_2}$

↗ can use any other norm

When stopping the iteration at $k = K$ set

$$\tilde{\lambda} \leftarrow x_{K+1}^* \cdot A \cdot x_{K+1}$$

Suppose that A is diagonalizable

$$A = S \cdot \Lambda \cdot S^{-1}$$

$(s_1 \dots s_n)$ ↗ $\text{diag}(\lambda_1, \dots, \lambda_n)$

s_i unit eigenvectors

with $|\lambda_1| > |\lambda_2| \geq \dots \geq |\lambda_n|$.

Note

$$x_0 = \alpha_1 s_1 + \dots + \alpha_n s_n$$

with $\alpha_1 \neq 0$

linear combination of the eigenvectors

Then

$$\begin{aligned} A^k x_0 &= \alpha_1 \lambda_1^k s_1 + \dots + \alpha_n \lambda_n^k s_n \\ &= \alpha_1 \lambda_1^k \left(s_1 + \sum_{j=2}^n \frac{\alpha_j}{\alpha_1} \left(\frac{\lambda_j}{\lambda_1} \right)^k s_j \right) \end{aligned}$$

This implies that for $x_k = \frac{A^k x_0}{\|A^k x_0\|_2}$
we have that

$$\|x_k - s_1\|_2 \leq O\left(\left|\frac{\lambda_2}{\lambda_1}\right|^k\right)$$

It can also be shown that for $\tilde{\lambda}_k = x_k^* \cdot A \cdot x_k$

$$|\tilde{\lambda}_k - \lambda_1| \leq O\left(\left|\frac{\lambda_2}{\lambda_1}\right|^k\right)$$

The number of correct digits in base b of this approximation

$$-\log_b \|x_k - s_1\|_2, -\log_b |\tilde{\lambda}_k - \lambda_1| \geq \log_b \left(\left|\frac{\lambda_1}{\lambda_2}\right|\right) \cdot k + \text{constant}$$

Minor drawback: the assumption $\boxed{\lambda_1 \neq 0}$.

If x_0 is random, this holds with very high probability.
Even if this is not the case, round offs during the iteration will ensure a non-zero component in the direction of s_1 . Many times we have a reasonable guess for x_0 .

In the PageRank algorithm: x_0 is the score vector of the previous web.

Major drawback: it converges to a pair eigenvalue/eigenvector only if there is dominant one, and the convergence is linear with ratio depending on $|\lambda_1/\lambda_2|$. This is not the case in

many situations:

- orthogonal matrices (all eigenvalues have absolute value 1)
- real matrices with complex (non real) eigenvalues.
(the eigenvalues come in pairs of conjugate complex numbers)

In the PageRank algorithm, $\lambda_1 = 1$ is the dominant eigenvalue and $\frac{|\lambda_2|}{|\lambda_1|} \leq 0.85$.

Advantage: the implementation only uses matrix-vector multiplications

In PageRank: the matrix-vector multiplication reduces to multiplication with the link matrix, which is sparse.

IX.2 Orthogonal iteration [D, § 4.4.3], [GVL, § 7.3.2]

This is a generalization of the power method, that can be used to compute higher dimensional invariant subspaces. Let $1 \leq r \leq n$.

Given Q_0 $n \times r$ unitary matrix

for $k = 0, 1, 2, \dots$

$$Y_{k+1} \leftarrow A \cdot Q_k$$

$$Q_{k+1} R_{k+1} = Y_{k+1}$$

QR factorization

Orthogonal iteration

Suppose again A diagonalizable

$$A = S \cdot \underset{\substack{\text{diag}(\lambda_1, \dots, \lambda_n)}}{\Lambda} \cdot S^{-1}$$

with $|\lambda_1| \geq \dots \geq |\lambda_r| > |\lambda_{r+1}| \geq \dots \geq |\lambda_n|$

We have that

$$\text{span}(Q_k) \stackrel{\substack{\text{linear space generated} \\ \text{by the columns of } Q_k}}{=} \text{span}(Y_k) = \dots = \text{span}(A^k Q_0)$$

$Y_k = Q_k \cdot R_k$

Hence Q_k gives an orthonormal basis of $\text{span}(A^k Q_0)$
prevents under/overflow and collapsing dimensions:
For $k \rightarrow \infty$ this should converge to an orthonormal
basis of the invariant subspace generated by
the dominant eigenvectors s_1, \dots, s_r :

$$\text{span}(Q_k) = \text{span}(A^k Q_0) \xrightarrow{k \rightarrow \infty} \text{span}(s_1, \dots, s_r)$$

For $r=1$ it coincides with the power method, because

$$x_{k+1} \| A y_{k+1} \|_2 = y_{k+1}$$

is the QR factorization of the vector y_{k+1} .

In this case $r=1$, $Q_0, Q_1, Q_2, \dots \in \mathbb{C}^{n \times 1}$

is the sequence of vectors produced by the power method.

Also for $1 \leq s < n$ the first s columns of Q_k coincide with the orthogonal iteration of rank s starting with the first s columns of Q_0 :

the QR factorization is compatible with restricting columns:

$$\begin{array}{c} \overbrace{\hspace{2cm}}^n \\ \begin{array}{|c|c|} \hline A' & A \\ \hline \end{array} \\ \begin{array}{c} m \\ n' \end{array} \end{array} = \begin{array}{c} \overbrace{\hspace{2cm}}^n \\ \begin{array}{|c|c|} \hline Q' & Q \\ \hline \end{array} \\ \begin{array}{c} m \\ n' \end{array} \end{array} \begin{array}{c} \begin{array}{|c|c|} \hline R' & R \\ \hline \end{array} \\ \begin{array}{c} n' \\ n \end{array} \end{array}$$

$A' = Q' R'$ is the QR factorization of A .

We can then set $r=n$, which would run the orthogonal iteration for all n (and initial matrix given by the first n columns of Q_0)

To study the convergence of the orthogonal iteration, suppose

$$|\lambda_1| > \dots > |\lambda_n|$$

Then for a "typical" Q_0 we have that for $k \rightarrow \infty$

$$Q_k \rightarrow Q \quad n \times n \text{ unitary}$$

$$Q_k^* \cdot A \cdot Q_k \rightarrow T \quad \text{upper triangular}$$

$$\text{s.t. } A = Q \cdot T \cdot Q^* \quad (\text{Schur decomposition})$$

see [D, Theorem 4.8]

IX.3 The QR iteration [D, §4.4.4] [GVL, §7.3.3]

Set

$$T_k := Q_k^* \cdot A \cdot Q_k \quad k=0, 1, 2, \dots$$

The QR iteration arises when considering how to compute T_k directly from its predecessor T_{k-1} .

We have

$$T_{k-1} = Q_{k-1}^* \cdot A \cdot Q_{k-1} \stackrel{\text{by definition of } Q_k}{=} (Q_{k-1}^* \cdot Q_k) \cdot R_k \quad (*)$$

$$T_k = Q_k^* \cdot A \cdot Q_k = (Q_k^* \cdot A \cdot Q_{k-1}) Q_{k-1}^* Q_k = R_k \cdot (Q_{k-1}^* \cdot Q_k)$$

$$Q_{k-1}^* \cdot A = R_k \cdot Q_k$$

We have that $Q_{k-1}^* \cdot Q_k$ are unitary and (*) is the QR factorization of T_{k-1}

Drawbacks:

- * A single QR iteration costs $O(n^3)$ flops.
- * Convergence (when it exists) is only linear.

We will next study how to overcome these practical difficulties.

[D, § 4.4.6], [GVL, § 7.4.1, 7.4.2 & 7.4.3]

IX. § The real Schur form and the Hessenberg reduction

Matrices from applications have real entries. Hence we concentrate on the real version of the QR iteration:
For $A \in \mathbb{R}^{n \times n}$ choose $Q_0 \in \mathbb{R}^{n \times n}$ orthogonal and set:

$$H_0 \leftarrow Q_0^T \cdot A \cdot Q_0$$

For $k = 1, 2, \dots$

$$H_{k-1} = Q_k \cdot R_k \quad (\text{QR factorization})$$

$$H_k \leftarrow R_k \cdot Q_k$$

If A has complex eigenvalues then H_k cannot converge to a triangular matrix. In this case we ~~not~~ content ourselves with convergence to the real Schur form:

$$A = Q \cdot T \cdot Q^T$$

orthogonal

block upper triangular

with $T = \begin{pmatrix} T_{11} & T_{12} & \dots & T_{1n} \\ 0 & \ddots & & \vdots \\ \vdots & & \ddots & T_{k-1,n} \\ 0 & \dots & 0 & T_{nn} \end{pmatrix}$

and $T_{ii} \begin{cases} 1 \times 1 \\ 2 \times 2 \text{ with complex conjugate eigenvalues} \end{cases}$

To implement the QR iteration, we have to choose Q_0 carefully st.

$$H_0 = Q_0^T \cdot A \cdot Q_0 \quad \text{is } \underline{\text{Hessenberg}} \\ (h_{ij} = 0 \text{ for } i > j+1)$$

In this way, the complexity of each iteration is lowered to $O(n^2)$ flops.

This Hessenberg reduction is a variation of the QR factorization, and can be done with a sequence of Householder reflections.

We illustrate it in the 5×5 case:

1) Choose $P_1 = \begin{pmatrix} 1 & & & & \\ & \tilde{P}_1 & & & \\ & & & & \end{pmatrix}$ st.
 $\xrightarrow{4 \times 5 \text{ Householder}}$

$$P_1 A = \begin{pmatrix} x & x & x & x & x \\ x & x & x & x & x \\ 0 & x & x & x & x \\ 0 & x & x & x & x \\ 0 & x & x & x & x \end{pmatrix} \quad \text{and} \quad A_1 = P_1 A P_1^T = \begin{pmatrix} x & x & x & x & x \\ x & x & x & x & x \\ 0 & x & x & x & x \\ 0 & x & x & x & x \\ 0 & x & x & x & x \end{pmatrix}$$

P_1 leaves the 1st row of $P_1 A$ unchanged

P_1^T leaves the 1st column of $(P_1 A) \cdot P_1^T$ unchanged, including the 0's.

2) Choose $P_2 = \begin{pmatrix} 1 & & & & \\ & \tilde{P}_2 & & & \\ & & & & \end{pmatrix}$ st.
 $\xrightarrow{3 \times 4 \text{ Householder}}$

$$P_2 A_1 = \begin{pmatrix} x & x & x & x & x \\ x & x & x & x & x \\ 0 & x & x & x & x \\ 0 & 0 & x & x & x \\ 0 & 0 & x & x & x \end{pmatrix} \quad \text{and} \quad A_2 = P_2 A_1 P_2^T = \begin{pmatrix} x & x & x & x & x \\ x & x & x & x & x \\ 0 & x & x & x & x \\ 0 & 0 & x & x & x \\ 0 & 0 & x & x & x \end{pmatrix}$$

P_2 leaves the 1st and 2nd rows of $P_2 \cdot A_1$ unchanged
 P_2^T leaves the 1st and 2nd columns of $(P_2 A_1) \cdot P_2^T$ unchanged

3) Choose $P_3 = \begin{pmatrix} I_3 & \\ & \tilde{P}_3 \end{pmatrix}$ st.

\tilde{P}_3 is a 2×2 Householder

$$P_3 \cdot A_2 = \begin{pmatrix} x & x & x & x & x \\ x & x & x & x & x \\ 0 & x & x & x & x \\ 0 & 0 & x & x & x \\ 0 & 0 & 0 & x & x \end{pmatrix} \quad \text{and} \quad P_3 \cdot A_2 \cdot P_3^T = \begin{pmatrix} x & x & x & x & x \\ x & x & x & x & x \\ 0 & x & x & x & x \\ 0 & 0 & x & x & x \\ 0 & 0 & 0 & x & x \end{pmatrix}$$

In general

$$(P_{n-2} \cdots P_1) \cdot A \cdot P_1^T \cdots P_{n-2}^T = H_0 \quad \text{Hessenberg}$$

$$\parallel$$

$$Q_0^T \cdot A \cdot Q_0$$

with $Q_0 = P_1 \cdots P_{n-2}$ $n \times n$ -orthogonal
 $(P_i \text{ is symmetric}).$

The complexity of computing Q_0 is $5n^3$ flops.

Now we turn to each QR step.

Important observation: the Hessenberg form is preserved during the QR iteration: Let

$$H = Q \cdot R \quad \text{and} \quad H_+ = R \cdot Q$$

Hessenberg

Since R is upper triangular, the j th column of Q is a linear combination of the first j columns of H :

$$H = (h_1 \dots h_n)$$

$$Q = (q_1 \dots q_n)$$

then

$$R = (r_{ij})$$

$$S = R^{-1} = (s_{ij})$$

$$q_j = \sum_{k=1}^j s_{kj} h_k$$

and so Q is Hessenberg.

Similarly the j th row of H_+ is a linear combination of the last j rows of Q and so H_+ is Hessenberg.

The QR factorization of H is with $n-1$ Givens rotations computed

$$Q^T \cdot H = R$$

upper triangular

$$\parallel$$

$$(G_{n-1}^T \dots G_1^T) \cdot H$$

with $G_i = \text{Givens}(i, i+1, \theta_i)$. Then

$$H_+ = R \cdot Q = R \cdot (G_1 \dots G_{n-1})$$

This requires $6n^2 + O(n)$ flops.

IX.5 Deflation

[GVL, § 7.5.4]

A Hessenberg matrix $H \in \mathbb{R}^{n \times n}$ is reduced if it has a zero subdiagonal entry. In this case

$$H = \begin{pmatrix} H_{11} & H_{12} \\ 0 & H_{22} \end{pmatrix} \begin{matrix} p \\ n-p \end{matrix} \quad \begin{matrix} 1 \leq p < n \\ p \quad n-p \end{matrix}$$

and the problem decouples into two smaller problems H_{11} and H_{22} (deflation). Typically occurs when $p = n-1$ or $p = n-2$.

In practice, this is done when a subdiagonal entry is sufficiently small, e.g. if

$$|h_{p+1,p}| \leq c \cdot \varepsilon \cdot (|h_{p,p}| + |h_{p+1,p+1}|)$$

for a small constant c . ε machine epsilon

IX.6 The shifted QR iteration [GVL, § 7.5.2 and § 7.5.3]

For $\mu \in \mathbb{R}$ consider the iteration

$$H \leftarrow Q_0^T A Q_0$$

Hessenberg reduction

For $k = 1, 2, \dots$

$$H - \mu \mathbb{1}_n = Q \cdot R$$

QR factorization

$$H \leftarrow R \cdot Q + \mu \mathbb{1}_n$$

Each H is Hessenberg and orthogonally similar to A :

$$\tilde{H} = R \cdot Q + \mu \mathbb{1}_n = Q^T (QR + \mu \mathbb{1}_n) Q = Q^T H \cdot Q$$

If we shift by an eigenvalue $\mu \in \lambda(A)$,
deflation occurs in one step:

[GvL, Theorem 7.5.1]

$$\tilde{h}_{n,n-1} = 0 \quad \text{and} \quad \tilde{h}_{nn} = \mu$$

In practice, we recognize that the QR iteration is
converging when

$\tilde{h}_{n,n-1}$ is small

Then we use $\mu = h_{nn}$, which gives a
quadratically converging algorithm: if

$$|\tilde{h}_{n,n-1}| \leq \eta \|H\| \quad \text{with } 0 \leq \eta \ll 1$$

then for $\mu = h_{nn}$ we have that

$$|\tilde{h}_{n,n-1}| \leq O(\eta^2 \|H\|)$$

the number of correct digits of $h_{nn} \approx \lambda \in \lambda(A)$
doubles at each step.

IX.7 The double shift strategy [GVL, § 7.5.4 & § 7.5.5]

The previous strategy ~~fails~~ when A has a complex (non real) eigenvalue $\lambda \in \mathbb{C} \setminus \mathbb{R}$.

In this case we might perform two single shifts QR steps in succession, using

λ and $\bar{\lambda}$ as shifts:

$$H - \lambda \mathbb{1}_n = Q_1 R_1$$

$$H_1 \leftarrow R_1 Q_1 + \lambda \mathbb{1}_n$$

$$H_1 - \bar{\lambda} \mathbb{1}_n = Q_2 R_2$$

$$H_2 \leftarrow R_2 Q_2 + \bar{\lambda} \mathbb{1}_n$$

Set $Q = Q_1 Q_2$, $R = R_2 R_1$, $M = QR$

Manipulating these equation it can be shown that

$$\begin{aligned} M &= (H - \lambda \mathbb{1}_n)(H - \bar{\lambda} \mathbb{1}_n) \\ &= H^2 - (\lambda + \bar{\lambda})H + \lambda \bar{\lambda} \mathbb{1}_n \end{aligned}$$

$\swarrow |\lambda|^2$
 $\nwarrow 2\operatorname{Re}(\lambda)$

and so M is real. It can also be shown that

$$H_2 = Q^T \cdot H \cdot Q \in \mathbb{R}^{n \times n}$$

To avoid complex arithmetics, we would like to pass from H to H_2 directly and using only real numbers.

The direct strategy would be:

- compute $M = H^2 - 2\operatorname{Re}(\lambda)H + |\lambda|^2 I_n$
- — the QR factorization $M = QR$
- set $H_2 = Q^T H Q$

The first step requires $O(n^3)$ and so it is not practical.

Fortunately, there is a strategy for computing this double shift in $O(n^2)$ flops: the implicit double shift, or Francis QR step:

- Compute $M \cdot e_1$ (1st column of M)
- Determine a Householder matrix P_0 s.t.
 $P_0 \cdot M e_1 = \text{multiple of } e_1$
- Compute Householder matrices P_1, \dots, P_{n-2} s.t. if
 $Z = P_0 P_1 \dots P_{n-2}$

then $Z \cdot H \cdot Z$ is upper Hessenberg and the 1st columns of Q and Z are equal.

This is based on the implicit QR thm: if

$$Q^T A Q = H \quad \text{and} \quad Z^T A Z = Q$$

are Hessenberg, H is unreduced and Q and Z have the same first column, then

$$G = D^{-1} H D \quad \text{with } D = \operatorname{diag}(\pm 1, \dots, \pm 1)$$

Hence if $Q^T H Q$ and $Z^T H Z$ are unreduced, they are essentially equal. Else we can decouple the problem into smaller unreduced subproblems.

In detail:

$$M e_1 = \begin{pmatrix} x \\ y \\ z \\ 0 \\ \vdots \\ 0 \end{pmatrix}$$

with

$$x = h_{11}^2 + h_{12} h_{21} - \overset{2 \operatorname{Re}(\lambda)}{s} h_{11} + \overset{|z|^2}{t}$$

$$y = h_{21} (h_{11} + h_{22} - s)$$

$$z = h_{21} h_{32}$$

The computation of $M e_1$ and P_0 takes $O(n)$ flops. A similarity with P_0 only changes rows and columns 1, 2 and 3:

$$P_0^T H P_0 = \begin{pmatrix} x & x & x & x & x & x \\ x & x & x & x & x & x \\ x & x & x & x & x & x \\ x & x & x & x & x & x \\ 0 & 0 & 0 & x & x & x \\ 0 & 0 & 0 & 0 & x & x \end{pmatrix}$$

The mission of P_1, \dots, P_{n-2} is to restore this matrix to Hessenberg form:

$$\begin{pmatrix} x & x & x & x & x & x \\ x & x & x & x & x & x \\ x & x & x & x & x & x \\ 0 & x & 0 & x & x & x \\ 0 & 0 & 0 & 0 & x & x \end{pmatrix} \xrightarrow{P_1} \begin{pmatrix} x & x & x & x & x & x \\ x & x & x & x & x & x \\ 0 & 0 & x & x & x & x \\ 0 & 0 & x & x & x & x \\ 0 & 0 & 0 & x & x & x \end{pmatrix} \xrightarrow{P_2} \begin{pmatrix} x & x & x & x & x & x \\ x & x & x & x & x & x \\ 0 & 0 & x & x & x & x \\ 0 & 0 & 0 & x & x & x \\ 0 & 0 & x & x & x & x \end{pmatrix}$$

etc

We have that

P_0 and $Z = P_0 P_1 \dots P_{n-2}$
have the same 1st column and so

$$Z e_1 = Q e_1$$

and so Z and Q coincide (up to signs).
To implement this strategy, we check when

$h_{n-1, n-2}$ is small

Then the eigenvalues of

$$\begin{pmatrix} h_{n-1, n-1} & h_{n-1, n} \\ h_{n, n-1} & h_{n, n} \end{pmatrix}$$

approximate eigenvalues of A . In this case we set

$$\begin{cases} s = h_{n-1, n-1} + h_{n, n} & (\text{trace}) \\ t = h_{n-1, n-1} \cdot h_{n, n} - h_{n-1, n} h_{n, n-1} & (\text{determinant}) \end{cases}$$

and gives quadratic convergence: on average,
two QR iterations are needed per eigenvalue,
and the overall process costs $O(n^3)$ flops